


# Chapter 5

# Propositions and Inference

**Textbook:** *Artificial Intelligence Foundations of Computational Agents*, David L. Poole and Alan K Mackworth, Cambridge University Press.



# Introduction

- This chapter presents several **reasoning formalisms** that use **propositions**.
- Here, a simple form of a **knowledge base (KB)** is presented, composed of **facts** and **rules** related to a problem world.
  - An **agent** can use such a **KB**, together with its **observations**, to determine the ***truth*** in the world.
- When **queried** about “*what must be **true** given a **KB**,*” it **answers** the **query** without enumerating all the possible worlds.

# Propositions

- **Statements** about the world provide **constraints** about *what could be **true***
  - **Constraints** could be specified extensionally as tables of legal assignments to **variables** or **formulas**.
- There are several reasons for using **propositions** for specifying **constraints** and **queries**:
  - It gives a more concise and readable logical statement about the **relationship** among **variables**.
  - The form of knowledge can be exploited to make **reasoning** more efficient.
    - They are **modular**, so small changes to the problem result in minor changes to the **KB**.

# Syntax of Propositional Calculus

- A **proposition** is a ***sentence*** written in a language with a **truth value** (either ***true*** or ***false***) in a world.
  - A **proposition** is built from **atomic propositions** using **logical connectives**.
  - We use the convention that **propositions** consist of *letters, digits,* and the *underscore* (‘\_’ ) and start with a **lowercase letter**.
  - For example, ***ai\_is\_fun***, ***live\_outside***, and ***sunny*** are all propositional atoms.
- **Propositions** can be built from more straightforward propositions using **logical connectives**
  - A proposition or logical formula is either an **atomic proposition** or a **compound proposition**

# Syntax of Propositional Calculus

- a **compound proposition** of the form

$\neg p$	“not $p$ ”	<b>negation</b> of $p$
$p \wedge q$	“ $p$ and $q$ ”	<b>conjunction</b> of $p$ and $q$
$p \vee q$	“ $p$ or $q$ ”	<b>disjunction</b> of $p$ and $q$
$p \rightarrow q$	“ $p$ implies $q$ ”	<b>implication</b> of $q$ from $p$
$p \leftarrow q$	“ $p$ if $q$ ”	<b>implication</b> of $p$ from $q$
$p \leftrightarrow q$	“ $p$ if and only if $q$ ”	<b>equivalence</b> of $p$ and $q$

where  $p$  and  $q$  are propositions.

The operators  $\neg, \wedge, \vee, \rightarrow, \leftarrow$  and  $\leftrightarrow$  are **logical connectives**.

# Syntax of Propositional Calculus

$p$	$q$	$\neg p$	$p \wedge q$	$p \vee q$	$p \leftarrow q$	$p \rightarrow q$	$p \leftrightarrow q$
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>false</i>
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>

Figure 5.1: Truth table defining  $\neg$ ,  $\wedge$ ,  $\vee$ ,  $\leftarrow$ ,  $\rightarrow$ , and  $\leftrightarrow$

# Semantics of the Propositional Calculus

- **Semantics** defines the **meaning** of the sentences of a language.
- When the **sentences** are about a world, **semantics** specifies ***how to put symbols of the language into correspondence with the world.***
- The **truth of atoms** gives the **truth** of other propositions in interpretations (see **Figure 5.1**).
- An interpretation consists of a **function  $\pi$**  that maps **atoms** to  **$\{true, false\}$** .
  - If  $\pi(a) = true$ , atom  **$a$**  is  **$true$**
  - If  $\pi(a) = false$ , atom  **$a$**  is  **$false$**
  - Suppose there are **three** atoms:  **$ai\_is\_fun$** ,  **$happy$** , and  **$light\_on$** ;  
 **$\pi(ai\_is\_fun) = true$ ,  $\pi(happy) = false$ ,  $\pi(light\_on) = true$**

# Semantics of the Propositional Calculus

**Example 5.20** Consider the knowledge base  $KB_2$ :

$$false \leftarrow a \wedge b.$$

$$a \leftarrow c.$$

$$b \leftarrow d.$$

$$b \leftarrow e.$$

Either  $c$  is false or  $d$  is false in every model of  $KB_2$ . If they were both true in some model  $I$  of  $KB_2$ , both  $a$  and  $b$  would be true in  $I$ , so the first clause would be false in  $I$ , a contradiction to  $I$  being a model of  $KB_2$ . Similarly, either  $c$  is false or  $e$  is false in every model of  $KB_2$ . Thus,

$$KB_2 \models \neg c \vee \neg d$$

$$KB_2 \models \neg c \vee \neg e.$$



# Semantics of the Propositional Calculus

- A **KB** is a set of **true** propositions.
- An element of the **KB** is an **axiom**.
- A **model** of a **KB** is an interpretation in which all the propositions in **KB** are **true**.
- If a knowledge base is **KB**, and **g** is a **proposition**, **KB**  $\models$  **g** is g logically follows from KB, or KB entails g.
- The logical entailment “**KB**  $\models$  **g**” is a semantic relation between a set of propositions in **KB** and an external proposition, **g**.
- Both **KB** and **g** are **symbolic** so that they can be represented in a computer.

# Propositional Constraints

- The class of **propositional satisfiability problems** has:
  - **Boolean variables**: If  $X$  is a Boolean variable,  $X = \text{false}$  as  $\neg x$ . Thus, given a Boolean variable *Happy*, the proposition *happy* means *Happy* = *true*, and  $\neg \text{happy}$  means *Happy* = *false*.
  - **Clausal constraints**: a **clause** is **disjoints of atoms** and is expressed as  $(I_1 \vee I_2 \vee \dots \vee I_k)$ , where each  $I_i$  is **literal**.
    - A **literal** is an **atom** or the **negation** of an **atom**
    - A **clause** is satisfied in a possible world if at least one of the **literals** that make up the clause is *true* in that possible world.

# Propositional Constraints

**Example 5.4** The clause  $happy \vee sad \vee \neg living$  is a constraint among the variables *Happy*, *Sad*, and *Living*, which is true if *Happy* has value *true*, *Sad* has value *true*, or *Living* has value *false*. The atoms *happy* and *sad* appear positively in the clause, and *living* appears negatively in the clause.

The assignment  $\neg happy, \neg sad, living$  violates the constraint of clause  $happy \vee sad \vee \neg living$ . It is the only assignment of these three variables that violates this clause.

# Propositional Constraints

- It is possible to convert any finite **CSP** into a **Propositional Satisfiable Problem (PSP)**:
  - A **CSP** variable  $Y$  with domain  $\{v_1, \dots, v_k\}$  can be converted into  $k$  **Boolean variables**  $\{Y_1, \dots, Y_k\}$ , where  $Y_i$  is **true** when  $Y$  has value  $v_i$  and is **false** otherwise. Each  $Y_i$  is called an **indicator variable**. Thus  $k$  atoms,  $y_1, \dots, y_k$ , are used to represent the **CSP variable**.
  - There are **constraints** that specify that  $y_i$  and  $y_j$  cannot both be **true** when  $(i \neq j)$ . There is a **constraint** that one of the  $y_i$  must be **true**. Thus, the **KB** contains the **clauses**:  $\neg y_i \vee \neg y_j$  for  $i < j$  and  $y_1 \vee \dots \vee y_k$ .
  - There is a **clause** for each **false** assignment in each **constraint**; for example, the clauses  $(a \vee b \vee c)$  and  $(a \vee b \vee \neg c)$  can be combined with  $(a \vee b)$ .

# Clausal Form for Consistency Algorithms

- **Consistency algorithms** can be made more efficient for **propositional satisfiability problems (PSP)**
  - When there are only *two values*, **pruning** a value from the domain is equivalent to assigning the opposite value
  - Thus, if ***X*** has domain ***{true, false}***, pruning ***true*** from the domain of ***X*** is the same as assigning ***X*** to have the value ***false***
- **Arc consistency** can be used to **prune** the set of values and the set of constraints.

# Clausal Form for Consistency Algorithms

- Assigning a value to a **Boolean variable** can simplify the **set of constraints**:
  - If  **$X$**  is assigned ***true***, all of the **clauses** with  **$X = \text{true}$**  become redundant; they are automatically satisfied. These clauses can be removed. Similarly, if  **$X$**  is assigned ***false***, clauses containing  **$X = \text{false}$**  can be removed.
  - If  **$X$**  is assigned ***true***, any **clause** with  **$X = \text{false}$**  can be simplified by removing  **$X = \text{false}$**  from the **clause**. Similarly, if  **$X$**  is assigned ***false***, then  **$X = \text{true}$**  can be removed from any **clause** it appears in. This step is called **unit resolution**.
  - After **pruning the clauses**, there is a **clause** that contains just one assignment,  **$Y = v$** , the other value can be removed from the domain of  **$Y$** . This is a form of **arc consistency**
    - *If all of the assignments are removed from a clause, the constraints are unsatisfiable.*

# Clausal Form for Consistency Algorithms

- If a **variable** has the same value in all remaining **clauses**, and the algorithm must only find **one model**, it can assign that value to that variable
  - For example, if variable ***Y*** only appears as ***Y = true*** (i.e.,  $\neg y$  is not in any clause), then ***Y*** can be assigned the value ***true***.
  - That assignment does not remove all of the models;
  - A variable that has only one value in all of the clauses is called a **pure literal**.

# Propositional Definite Clauses

- The syntax of **definite propositional clauses** is defined as follows:
  - An atomic proposition or atom is the same as in propositional calculus.
  - A **definite clause** is of the form:  $a \leftarrow b_1 \wedge \dots \wedge b_m$ , where  **$a$  is the head of the clause and is always positive**, and  **$a$**  and each  **$b_i$**  are called **atoms**. It can be read as “ **$a$  if  $b_1$  and  $b_2$ . . . and  $b_m$** ”
    - $(a \leftarrow (b_1 \wedge b_2 \wedge \dots \wedge b_m)) = (a \vee \neg b_1 \vee \neg b_2 \dots \vee \neg b_m)$ 
      - If  **$m > 0$** , the clause is called a **rule**, where  $b_1 \wedge \dots \wedge b_m$  is the **body** of the clause ( “**if**” part of the rule).
      - If  **$m = 0$** , the **arrow can be omitted**, and the clause is called an **atomic clause** or a **fact** (it is the **clause with an empty body**).
  - **A knowledge base (KB)** is a set of **definite clauses**.



# Propositional Definite Clauses

**Example 5.6** The elements of the knowledge base in Example 5.20 are all definite clauses.

The following are *not* definite clauses:

$\neg \text{apple\_is\_eaten.}$

$\text{apple\_is\_eaten} \wedge \text{bird\_eats\_apple.}$

$\text{sam\_is\_in\_room} \wedge \text{night\_time} \leftarrow \text{switch\_1\_is\_up.}$

$\text{Apple\_is\_eaten} \leftarrow \text{Bird\_eats\_apple.}$

$\text{happy} \vee \text{sad} \vee \neg \text{alive.}$

The fourth statement is not a definite clause because an atom must start with a lower-case letter.

# Propositional Definite Clauses

- Note that a **definite clause** is a restricted form of a **clause**. For example, the **definite clause**  $(a \leftarrow b \wedge c \wedge d)$  is **equivalent** to the **clause**  $(a \vee \neg b \vee \neg c \vee \neg d)$ . That is,  $(a \leftarrow b \wedge c \wedge d) = (a \vee \neg b \vee \neg c \vee \neg d)$
- In general, a **definite clause** is equivalent to a clause with precisely one positive literal.

# Propositional Definite Clauses

- **Example 5.7** Consider how to axiomatize the electrical environment of **Figure 5.2** following the methodology for the user's view of semantics.

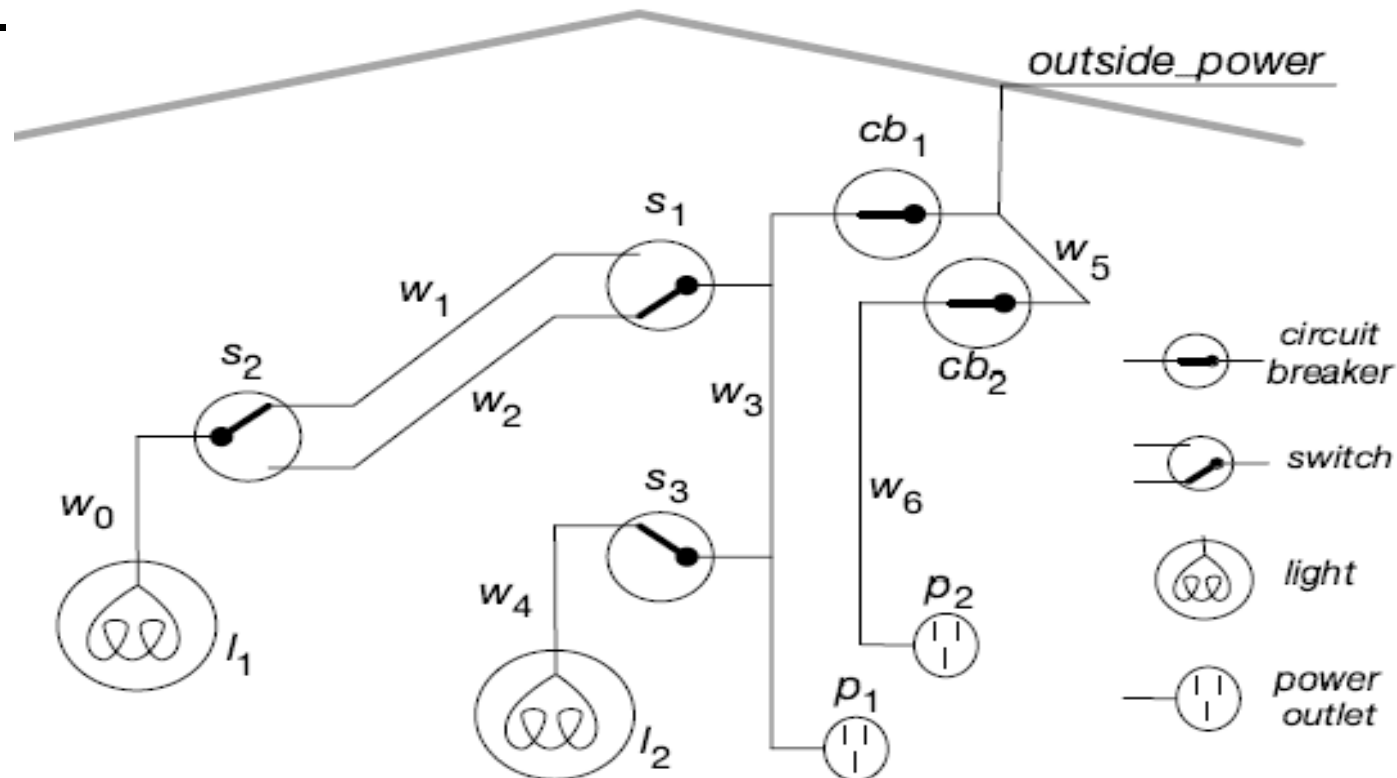


Figure 5.2: An electrical environment with components named

# Propositional Definite Clauses

- The designer may look at part of the domain and know that **light**  $l_1$  is **live** if **wire**  $w_0$  is **live** because they are connected but may not know whether  $w_0$  is **live**. Such knowledge is expressible in terms of rules.
- The **KB** consists of all of the **definite clauses**, whether specified as background knowledge or as observations.

$live\_l_1 \leftarrow live\_w_0.$   
 $live\_w_0 \leftarrow live\_w_1 \wedge up\_s_2.$   
 $live\_w_0 \leftarrow live\_w_2 \wedge down\_s_2.$   
 $live\_w_1 \leftarrow live\_w_3 \wedge up\_s_1.$   
 $live\_w_2 \leftarrow live\_w_3 \wedge down\_s_1.$   
 $live\_l_2 \leftarrow live\_w_4.$   
 $live\_w_4 \leftarrow live\_w_3 \wedge up\_s_3.$   
 $live\_p_1 \leftarrow live\_w_3.$   
 $live\_w_3 \leftarrow live\_w_5 \wedge ok\_cb_1.$   
 $live\_p_2 \leftarrow live\_w_6.$   
 $live\_w_6 \leftarrow live\_w_5 \wedge ok\_cb_2.$   
 $live\_w_5 \leftarrow live\_outside.$   
 $lit\_l_1 \leftarrow light\_l_1 \wedge live\_l_1 \wedge ok\_l_1.$   
 $lit\_l_2 \leftarrow light\_l_2 \wedge live\_l_2 \wedge ok\_l_2.$

# Questions and Answers

- A **query** is a way of asking whether a proposition is a **logical consequence** of a **KB**.
- A **query** is a question with the answer “**yes**” if the body is a logical **consequence** of the **KB** or
- the answer “**no**” if the body is **not a consequence** of the **KB**.

# Questions and Answers

**Example 5.8** Once the computer has been told the knowledge base of Example 5.7 (page 183), it can answer queries such as

ask *light*  $l_1$ .

for which the answer is *yes*. The query

ask *light*  $l_6$ .

has answer *no*. The computer does not have enough information to know whether or not  $l_6$  is a light. The query

ask *lit*  $l_2$ .

has answer *yes*. This atom is true in all models.

The user can interpret this answer with respect to the intended interpretation.

# Proofs

- A **theorem** is a **provable proposition**.
  - A **proof** is a mechanically derivable demonstration that a proposition logically follows from a **KB**.
- A **proof procedure** is a – possibly **non-deterministic algorithm** for deriving consequences of a **KB**
  - A **proof procedure** is **sound** concerning semantics if everything that can be derived from a **KB** is a **logical consequence** of the **KB**.
  - A **proof procedure** is **complete** concerning semantics if there is **proof of each logical consequence** of the **KB**.
  - **Two ways to construct proofs** for definite propositional clauses:
    - a **bottom-up procedure** and
    - a **top-down procedure**.

# Bottom-Up Proof Procedure

- A **bottom-up proof procedure** can be used to derive **all logical consequences of a KB**
  - The **bottom-up proof procedure** builds on atoms that have already been established.
  - We say that a **bottom-up procedure** is **forward chaining** on the **definite clauses**, in the sense of ***going forward from what is known*** rather than going backward from the query.
  - The general idea is based on **one rule of derivation** is a generalized form of the rule of **inference** called **modus ponens**:
    - If “ $h \leftarrow a_1 \wedge \dots \wedge a_m$ ” is a **definite clause** in the **KB**, and each  $a_i$  has been derived, then  $h$  can be derived.
- **Figure 5.3** gives a procedure for computing the **consequence set C** of a set KB of **definite clauses**.



# Bottom-Up Proof Procedure

- **Modus Ponens:** Create a conclusion from a rule and a fact. For example, consider the following **rule** and a **fact**:

Rule:  $(a \leftarrow b)$

Fact: **b**

Conclusion: **a**

# Bottom-Up Proof Procedure

---

```
1: procedure Prove_DC_BU(KB)
2:   Inputs
3:     KB: a set of definite clauses
4:   Output
5:     Set of all atoms that are logical consequences of KB
6:   Local
7:     C is a set of atoms
8:     C := {}
9:   repeat
10:    select " $h \leftarrow a_1 \wedge \dots \wedge a_m$ " in KB where  $a_i \in C$  for all  $i$ , and  $h \notin C$ 
11:    C :=  $C \cup \{h\}$ 
12:  until no more definite clauses can be selected
13:  return C
```

Figure 5.3: Bottom-up proof procedure for computing consequences of *KB*

# Bottom-Up Proof Procedure

**Example 5.9** Suppose the system is given the knowledge base  $KB$ :

$$a \leftarrow b \wedge c.$$

$$b \leftarrow d \wedge e.$$

$$b \leftarrow g \wedge e.$$

$$c \leftarrow e.$$

$$d.$$

$$e.$$

$$f \leftarrow a \wedge g.$$

One trace of the value assigned to  $C$  in the bottom-up procedure is

$$\{\}$$

$$\{d\}$$

$$\{e, d\}$$

$$\{c, e, d\}$$

$$\{b, c, e, d\}$$

$$\{a, b, c, e, d\}.$$

The algorithm terminates with  $C = \{a, b, c, e, d\}$ .

Thus,  $KB \vdash a$ ,  $KB \vdash b$ , and so on.

The last rule in  $KB$  is never used. **The bottom-up proof procedure** cannot derive  **$f$**  or  **$g$** . This is as it should be because there is a model of the  $KB$  in which  **$f$**  and  **$g$**  are both **false**.

# Top-Down Proof Procedure

- An alternative **proof method** is to search **backward** or **top-down** from a **query** to determine whether it is a logical consequence of the given **definite clauses**.
- **Top-down procedure** is also called **backward chaining**
- The **top-down proof procedure** can be understood in the **answer clause** (it is a definite clause with the head “**yes**”).
  - An **answer clause** is of the form: **yes**  $\leftarrow a_1 \wedge a_2 \wedge \dots \wedge a_m$  where **yes** is a special atom.
  - Intuitively, **yes** is going to be **true** exactly when the answer to the query is “**yes**.”
- If the **query** is (**ask**  $q_1 \wedge \dots \wedge q_m$ ), then the initial **answer clause** is **yes**  $\leftarrow q_1 \wedge \dots \wedge q_m$

# Top-Down Proof Procedure

- Given an **answer clause**, the **top-down algorithm** selects an atom in the body of the **answer clause**.
- Suppose it selects  $a_1$ . The atom selected is called a **subgoal**.
- The algorithm proceeds by doing steps of **resolution**
- In one step of **resolution**, it chooses a **definite clause** in **KB** with  $a_1$  as the **head**. If there is no such **clause**, the algorithm fails.
- The **resolvent** of the above **answer clause** on the selection  $a_1$  with the **definite clause** ( $a_1 \leftarrow b_1 \wedge \dots \wedge b_p$ ) is the **answer clause**:  
(**yes**  $\leftarrow b_1 \wedge \dots \wedge b_p \wedge a_2 \wedge \dots \wedge a_m$ )
- That is, the **subgoal** in the **answer clause** is replaced by the body of the chosen **definite clause**.
- An **answer** is an **answer clause** with an **empty body** ( $m = 0$ ), it is the **answer clause** (**yes**  $\leftarrow$ ) .

# Top-Down Proof Procedure

- **Figure 5.4** specifies a **non-deterministic procedure** for **solving a query**. It follows the definition of a derivation.
- In this procedure, ***G*** is the set of **atoms** in the body of the **answer clause**.
- The procedure is **nondeterministic**: **line 12** has to choose a **definite clause** to resolve against.
- If there are choices that result in ***G*** being the ***empty set***, the algorithm returns **yes**; otherwise, it ***fails***, and the answer is ***no***.
- This algorithm treats the body of a clause as a set of atoms, and ***G*** is also a set of atoms.
- An alternative is to have ***G*** as an ordered list of atoms, perhaps with an atom appearing more than once.

# Top-Down Proof Procedure

```
1: non-deterministic procedure Prove_DC_TD(KB, Query)
2:   Inputs
3:     KB: a set of definite clauses
4:     Query: a set of atoms to prove
5:   Output
6:     yes if  $KB \models Query$  and the procedure fails otherwise
7:   Local
8:     G is a set of atoms
9:      $G := Query$ 
10:  repeat
11:    select an atom a in G
12:    choose definite clause " $a \leftarrow B$ " in KB with a as head
13:     $G := B \cup (G \setminus \{a\})$ 
14:  until  $G = \{\}$ 
15:  return yes
```

Figure 5.4: Top-down definite clause proof procedure

# Top-Down Proof Procedure

**Example 5.10** Suppose the system is given the knowledge base:

$a \leftarrow b \wedge c.$

$b \leftarrow d \wedge e.$

$b \leftarrow g \wedge e.$

$c \leftarrow e.$

$d.$

$e.$

$f \leftarrow a \wedge g.$

It is asked the query: ask  $a.$

The following shows a derivation that corresponds to a sequence of assignments to  $G$  in the repeat loop of [Figure 5.4](#). Here we have written  $G$  in the form of an answer clause, and always selected the leftmost atom in the body:



# Top-Down Proof Procedure

$$\begin{aligned} \text{yes} &\leftarrow a \\ \text{yes} &\leftarrow b \wedge c \\ \text{yes} &\leftarrow d \wedge e \wedge c \\ \text{yes} &\leftarrow e \wedge c \\ \text{yes} &\leftarrow c \\ \text{yes} &\leftarrow e \\ \text{yes} &\leftarrow \end{aligned}$$

.

The following shows a sequence of choices, where the second definite clause for  $b$  was chosen. This choice does not lead to a proof.

$$\begin{aligned} \text{yes} &\leftarrow a \\ \text{yes} &\leftarrow b \wedge c \\ \text{yes} &\leftarrow g \wedge e \wedge c \end{aligned}$$

If  $g$  is selected, there are no rules that can be chosen. This proof attempt is said to fail.

.

# Top-Down Proof Procedure

- The **non-deterministic top-down** algorithm of **Figure 5.4**, together with a selection strategy, induces a **search graph**.
- Each **node** in the search graph represents an **answer clause**.
- The neighbors of a **node** (**yes**  $\leftarrow \mathbf{a}_1 \wedge \dots \wedge \mathbf{a}_m$ ), where  $\mathbf{a}_1$  is the **selected atom**, represent all of the possible answer clauses obtained by resolving on  $\mathbf{a}_1$ .
- There is a neighbor for each **definite clause** whose head is  $\mathbf{a}_1$ .
- The **goal nodes** of the search are of the form **yes**  $\leftarrow$  .

# Top-Down Proof Procedure

**Example 5.11** Given the knowledge base

$a \leftarrow b \wedge c.$	$a \leftarrow g.$	$a \leftarrow h.$
$b \leftarrow j.$	$b \leftarrow k.$	$d \leftarrow m.$
$d \leftarrow p.$	$f \leftarrow m.$	$f \leftarrow p.$
$g \leftarrow m.$	$g \leftarrow f.$	$k \leftarrow m.$
$h \leftarrow m.$	$p.$	

and the query

$\text{ask } a \wedge d.$

the search graph for an SLD derivation, assuming the leftmost atom is selected in each answer clause, is shown in [Figure 5.5](#).

# Top-Down Proof Procedure

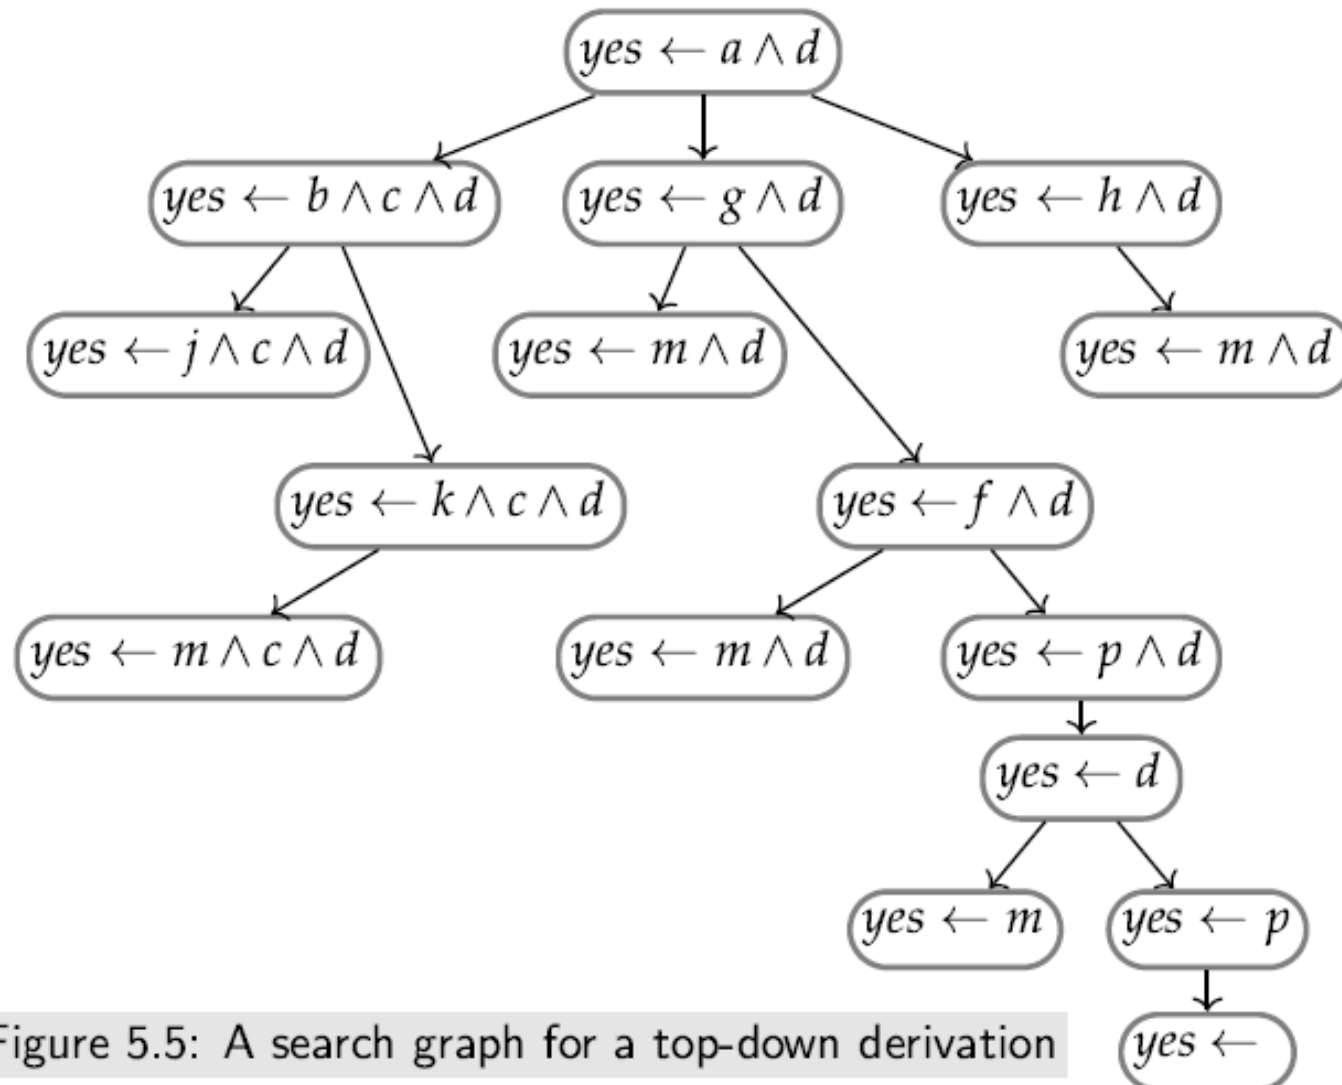


Figure 5.5: A search graph for a top-down derivation

# Top-Down Proof Procedure

- It is possible that the proof procedure can get into an **infinite loop**, as in the following example (without cycle pruning).

---

**Example 5.12** Consider the knowledge base and query:

$g \leftarrow a.$

$a \leftarrow b.$

$b \leftarrow a.$

$g \leftarrow c.$

$c.$

ask  $g.$

Atoms  $g$  and  $c$  are the only atomic logical consequences of this knowledge base, and the bottom-up proof procedure will halt with fixed point  $\{c, g\}$ . However, the top-down proof procedure with a depth-first search will go on indefinitely, and not halt if the first clause for  $g$  is chosen, and there is no cycle pruning.

# Background Knowledge and Observations

- Assume an **observation** is a set of **atomic propositions** which are **implicitly conjoined**.
- **Observations do not provide rules directly.**
- **Observation** is like **background knowledge** in a **KB**, allowing the agent to do something useful.
- At **design time** or **offline**, the **information arrives online** as **observations** from *users, sensors, and external knowledge sources*
  - For example, a **medical diagnosis program** may have knowledge represented as **definite clauses** about the possible diseases and symptoms, but *it would not know the actual symptoms manifested by a particular patient*
  - A simple way to acquire information from a user is to incorporate an **ask-the-user** mechanism into the **top-down proof procedure**.

# Knowledge-Level Explanation

- The explicit use of semantics allows explanation and debugging at the **knowledge level**
- To make a system usable by people, the system cannot just give an answer and expect the user to believe it
  - Consider the case of a system **advising doctors** who are legally responsible for the treatment they carry out based on the diagnosis.
    - The doctors must be convinced that the diagnosis is appropriate.
    - The system must be able to justify that its answer is correct.
    - The exact mechanism can explain how the system found a result and debugged the **KB**.

# Knowledge-Level Explanation

- **Three complementary** means of interrogation are used to explain the **relevant knowledge in a KB**:
  - (1) *how the question* is used to describe *how an answer was proved*,
  - (2) *why the question* is used to ask the system *why it is asking the user a question*, and
  - (3) *why not question* is used to ask *why an atom was not proven*.
- To explain how an answer was proved, a “**how**” question can be asked by a user when the system has returned the answer
  - The system provides the **definite clause** used to deduce the answer.
  - The user can ask “**why**” when asked a question.
  - The system **replies** by giving the rule that produced the question.
  - The user can then ask **why** the head of that rule was being proved.



# Knowledge-Level Explanation

- **How a System can Prove an Atom?**
- The first explanation procedure allows the user to ask “**how**” an atom was proved.
- If there is a proof for  **$g$** , either  **$g$**  must be an **atomic clause** or there must be a rule:  
$$g \leftarrow (a_1 \wedge \dots \wedge a_k)$$
 such that each  **$a_i$**  has been proved.

# Proving by Contradiction

- **Definite clauses** can be used in a proof by **contradiction** by allowing rules that give **contradictions**.
  - The **definite clause language** does not allow a contradiction to be stated.
- An **integrity constraint** is a clause of the form: **false**  $\leftarrow a_1 \wedge \dots \wedge a_k$ , where the  $a_i$  are atoms and **false** is a particular atom that is **false** in all interpretations.
- A **Horn clause** is either a **definite clause** or an **integrity constraint**.
- That is, a **Horn clause** has either **false** or a **normal atom** as its **head**.

# Proving by Contradiction

- **Integrity constraints** allow the system to prove that some **conjunction of atoms is *false*** in all models of a **KB**.
- $\neg p$  is the **negation of  $p$** , which is ***true*** in an interpretation when  $p$  is ***false*** in that interpretation,
- and  $p \vee q$  is the **disjunction** of  $p$  and  $q$ , which is ***true*** in an interpretation if  $p$  is ***true*** or  $q$  is ***true***, or both are ***true*** in the interpretation.
- The **integrity constraint:  $false \leftarrow a_1 \wedge \dots \wedge a_k$**  is logically equivalent to  $\neg a_1 \vee \dots \vee \neg a_k$ .
- A **Horn clause KB** can imply **negations of atoms**, as shown in **Example 5.19**.

# Proving by Contradiction

**Example 5.19** Consider the knowledge base  $KB_1$ :

$$false \leftarrow a \wedge b.$$

$$a \leftarrow c.$$

$$b \leftarrow c.$$

The atom  $c$  is false in all models of  $KB_1$ . To see this, suppose instead that  $c$  is true in model  $I$  of  $KB_1$ . Then  $a$  and  $b$  would both be true in  $I$  (otherwise  $I$  would not be a model of  $KB_1$ ). Because  $false$  is false in  $I$  and  $a$  and  $b$  are true in  $I$ , the first clause is false in  $I$ , a contradiction to  $I$  being a model of  $KB_1$ . Thus  $\neg c$  is true in all models of  $KB_1$ , which can be written as

$$KB_1 \models \neg c$$

# Proving by Contradiction

- Although the language of **Horn clauses** does not allow **disjunctions** and **negations** to be input, disjunctions of negations of atoms can be derived, as the following example shows.

**Example 5.20** Consider the knowledge base  $KB_2$ :

$$false \leftarrow a \wedge b.$$

$$a \leftarrow c.$$

$$b \leftarrow d.$$

$$b \leftarrow e.$$

Either  $c$  is false or  $d$  is false in every model of  $KB_2$ . If they were both true in some model  $I$  of  $KB_2$ , both  $a$  and  $b$  would be true in  $I$ , so the first clause would be false in  $I$ , a contradiction to  $I$  being a model of  $KB_2$ . Similarly, either  $c$  is false or  $e$  is false in every model of  $KB_2$ . Thus,

$$KB_2 \models \neg c \vee \neg d$$

$$KB_2 \models \neg c \vee \neg e.$$

# Proving by Contradiction

- A set of clauses is **unsatisfiable** if it has **no models**.
- A set of clauses is provably **inconsistent** concerning a **proof procedure**.
- If a proof procedure is **sound** and **complete**, a set of clauses is provably **inconsistent** if and only if it is **unsatisfiable**.
- It is always possible to find **a model** for a set of **definite clauses**.
- The interpretation with all atoms **true** is a model of any set of **definite clauses**.
- Thus, **a definite-clause KB is always satisfiable**.
- However, a set of **Horn clauses can be unsatisfiable**.

# Abduction

- **Abduction** is a form of reasoning where ***assumptions are made to explain observations.***
  - For example, *if an agent observes that some light is not working, it hypothesizes what is happening in the world to explain why the light is not working.*
- In **abduction**, an agent **hypothesizes** what may be ***true*** about an observed case.
- An agent determines what implies its observations – what could be ***true*** to make them ***true***.
- To formalize **abduction**, we use **Horn clauses** and **assumable**. The system is given:
  - a **KB**, a set of **Horn clauses**, and a **set A of atoms**, called the **assumable**, are the building blocks of **hypotheses** (see **Example 5.33**).

# Abduction

**Example 5.33** Consider the following simplistic knowledge base and assumables for a diagnostic assistant:

*bronchitis*  $\leftarrow$  *influenza*.

*bronchitis*  $\leftarrow$  *smokes*.

*coughing*  $\leftarrow$  *bronchitis*.

*wheezing*  $\leftarrow$  *bronchitis*.

*fever*  $\leftarrow$  *influenza*.

*fever*  $\leftarrow$  *infection*.

*soreThroat*  $\leftarrow$  *influenza*.

*false*  $\leftarrow$  *smokes*  $\wedge$  *nonsmoker*.

assumable *smokes, nonsmoker, influenza, infection*.



# Abduction

- If the agent observes **wheezing**, there are two minimal explanations: {**influenza**} and {**smokes**}
- These explanations imply **bronchitis** and **coughing**.
- If (**wheezing**  $\wedge$  **fever**) is observed, the **minimal explanations** are {**influenza**} and {**smokes**, **infection**}.
- If (**wheezing**  $\wedge$  **nonsmoker**) was observed, there is one **minimal explanation**: {**influenza**, **nonsmoker**}.
- The other explanation of **wheezing** is inconsistent with being a **non-smoker**.

# Abduction

**Example 5.34** Consider the knowledge base:

$alarm \leftarrow tampering.$

$alarm \leftarrow fire.$

$smoke \leftarrow fire.$

If *alarm* is observed, there are two minimal explanations:

$\{tampering\}$  and  $\{fire\}$ .

If  $alarm \wedge smoke$  is observed, there is one minimal explanation:

$\{fire\}$ .

Notice how, when *smoke* is observed, there is no need to hypothesize *tampering* to explain *alarm*; it has been **explained away** by *fire*.