

# An Introduction To Artificial Neural Network

## Reference(s):

***An Introduction to Neural Networks for Beginners***, Dr. Andy Thomas  
(Adventures in Machine Learning).

**<https://www.analyticsvidhya.com/blog/2021/11/a-comprehensive-guide-to-linear-regression-with-pytorch/>**

***Artificial Intelligence: A Guide to Intelligence Systems***, Michael Negnevitsky,  
Addison Wesley

# Introduction to Neural Networks

- An **artificial neural network (ANN)** or **neural network (NN)** is a software implementation of the neuronal structure of a human brain.
- The brain contains **neurons** which are kind of like biological switches.
  - These can change their **output state** depending on the strength of their electrical or chemical **input**.
  - The **NN** in a person's brain is a hugely interconnected network of **neurons**, where the **output** of any given neuron may be the **input** to thousands of other neurons (massively parallel structure!).

# Introduction to Neural Networks

- **Neural learning** occurs by repeatedly activating certain neural connections over others, reinforcing those connections.
- This makes them more likely to produce the desired outcome given a specified input
  - This learning involves **feedback** – when the **desired outcome** occurs, the neural connections causing that outcome become strengthened.

# Introduction to Neural Networks

- **NNs** attempt to simplify and mimic brain behavior.
- They can be trained in a ***supervised*** or ***unsupervised*** learning manner.
- In a **supervised NN**, the network is trained by providing matched **input-output data samples**, with the intention of getting the ANN to provide the **desired output** for a given **input**.

# Supervised NN An Example

- Consider an e-mail **spam filter** – the **input training data** could be the count of various words in the body of the email, and the **output training data** would be a classification of whether the e-mail was truly spam or not.
- If many examples of e-mails are passed through the NN allows the network to learn what input data makes it likely that an e-mail is a spam or not.
- This learning takes place by adjusting the **weights** of the NN connections.

# Unsupervised NN

- In **unsupervised learning**, the NN tries to “understand” and generate the **output** structure of the provided input data set “on its own” without its output pattern.
  - There is no supervised learning happening at this point

# Structure of an Artificial Neuron

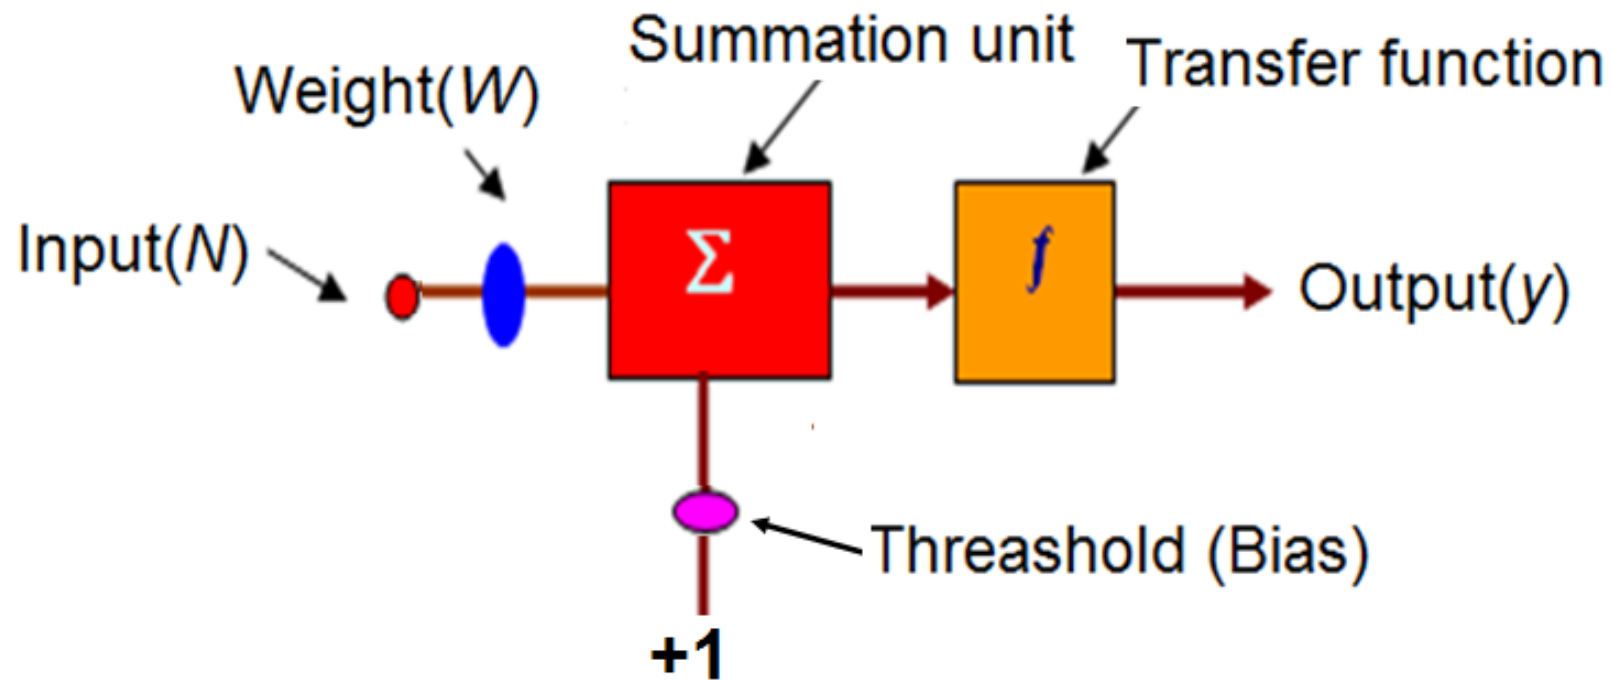
- The **neurons** that we are going to see here are not biological but are **Artificial Neurons**
- The **artificial Neurons** are extremely simple abstractions of **biological neurons**, realized as elements in *a program* or perhaps as *a circuit* made of silicon.
- Networks of these **artificial neurons** do not have a fraction of the power of the human brain, but they can merely be trained to perform useful functions.

# The Structure of an Artificial Neuron

- A *single-input, single-output* artificial neuron is shown in **figure 1**.
- The scalar input  **$N$**  is multiplied by the scalar weight  **$W$**  to form  **$W*N$** , one of the terms that are sent to the **summer unit**:
  - Where  **$N = \{n_1, n_2, n_3, \dots\}$**  and  **$W = \{w_1, w_2, w_3, \dots\}$**
  - The summer output is often referred to as the **net input**. It incorporates a **threshold** (or **bias**), the **weight of the +1 bias element** that goes into a **transfer function (or activation function)  $f$** , which determines and produces the scalar neuron output  **$y$** .



# The Structure of an Artificial Neuron



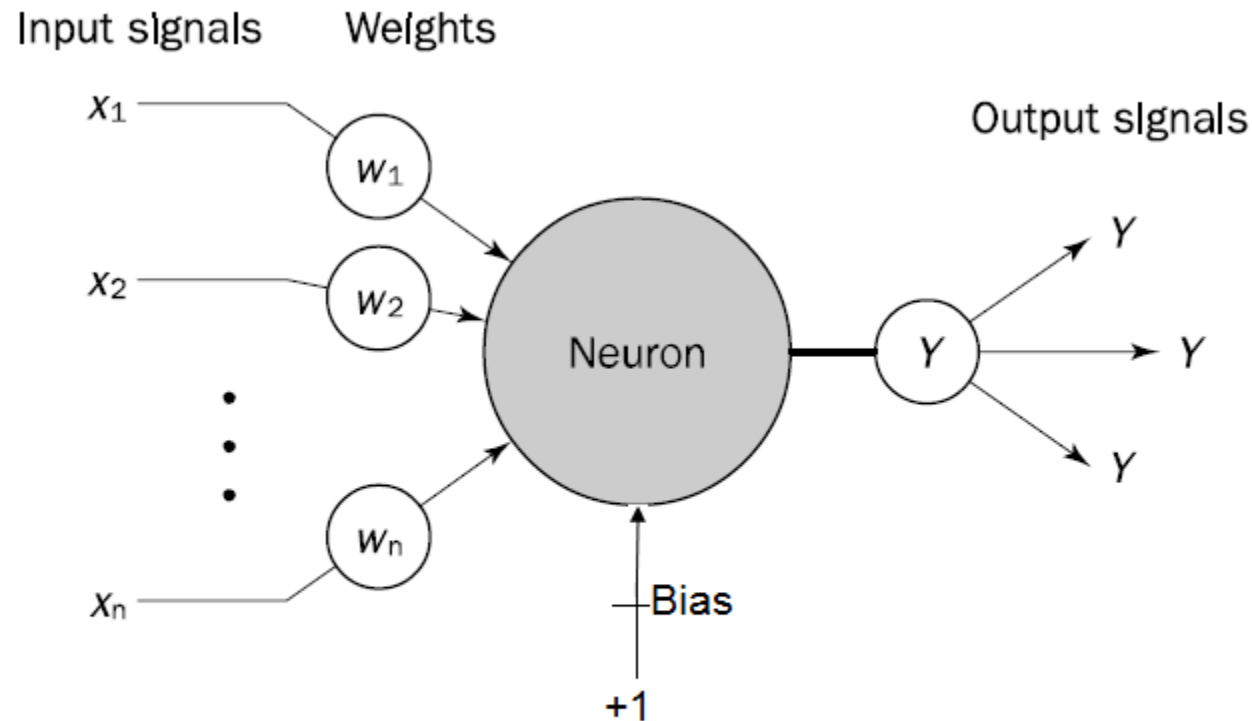
**Figure 1:** A single input, single output neuron

# Neuron as Simple Computing Element

- An artificial **neuron** receives several signals from its **input** links, computes with its **activation function** (**activation level**), and sends the result as an **output** through its **output links** (**figure 2** shows a neuron with  $n$  inputs)
- An activation function simulates the neuron in an ANN.
- The **input** signal can be *raw data* or *outputs* of other neurons
- The **output** signal can be either a *final solution* to a problem or an *input* to other neurons.

# Neuron as a Simple Computing Element

- The single neuron node which is shown in **figure 2** is called as a **perceptron**.



**Figure 2.** Single neuron with  $n$  inputs

# How Does a Neuron Determine its Output?

- The neuron computes the **weighted sum** of the input signals and compares the result with a **bias** value  $b$ .
- *If the **net input** is less than the  $b$ , then the neuron output is  $-1$*
- *But if the **net input** is greater than or equal to the **bias**  $b$ , the neuron becomes activated, and its output attains a value  $+1$  (McCulloch, 1943)*
- From **figure 2**, the **weighted\_sum**  $X$  can be given as:

$$\mathbf{X} = \sum_{i=1}^n (x_i w_i) \quad (1)$$

$$\text{where output, } \mathbf{Y} = \begin{cases} +1 & \text{if } X \geq b \\ -1 & \text{if } X < b \end{cases}$$

- Where  $x_i$  is the value of the input,  $w_i$  is the weight of input, and  $n$  is the number of inputs of the neuron.

# How Does a Neuron Determine its Output?

- The output of the neuron **Y** (from **Figure 2**) is estimated as;

$$Y = \begin{cases} +1 & \text{if } X \geq b \\ -1 & \text{if } X < b \end{cases}$$

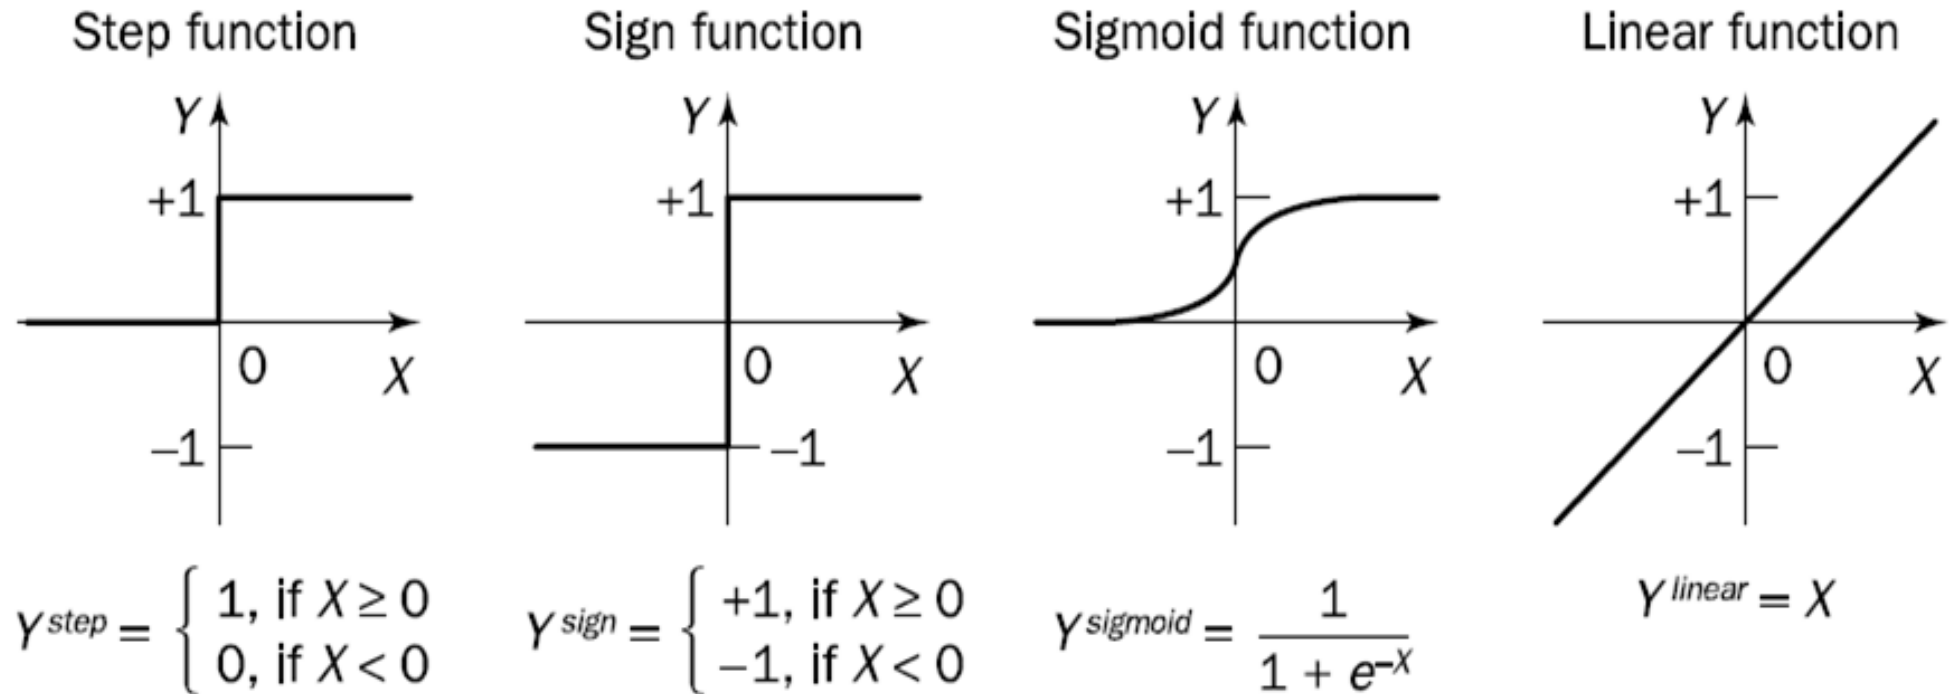
- This type of activation function is called a **sign activation** function
- Thus the **actual output** of the neuron with a **sign activation** function can be represented as:

$$Y = \mathbf{sign}[X] \quad ; \text{Where } X = \sum_{i=1}^n (x_i w_i) - b \quad \text{or} \quad \sum_{i=1}^n (x_i w_i) + b \quad (2)$$

Then output can be shown as,  $Y = \begin{cases} +1 & \text{if } X \geq 0 \\ -1 & \text{if } X < 0 \end{cases}$

- There are many activation functions: *step, sign, linear, sigmoid, etc.*
- Figure 3** shows the common activation functions of a neuron.

# Activation Functions of a Neuron



**Figure 3:** Activation functions of a neuron

# Activation Functions of a Neuron

- The **step** and **sign** activation functions are also called **hard limit** functions, are often used in **decision making neurons** for **classification** application
- The **sigmoid** function ( $Y_{sigmoid} = \frac{1}{1 + e^{-X}}$ , where **e** = 2.7183)

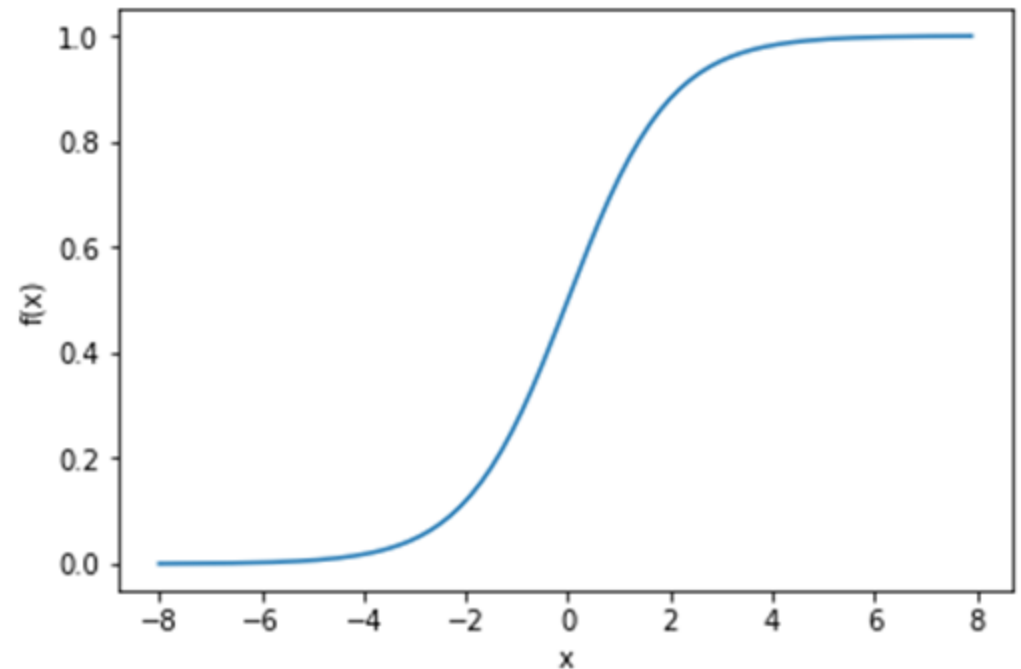
transforms the **input**, which can have any value between **plus** and **minus infinity**, into a reasonable value in the range between **0** to **1**

- Neurons with this function are used in the **back-propagation** networks
- The **linear function** ( $Y_{linear} = X$ ) provides an *output equal to the neuron weighted input*
  - Neurons with *linear function* are often used for *linear approximation*

# The Sigmoid Function

- The **Sigmoid** activation function: ( $Y_{sigmoid} = 1/(1 + e^{-x})$ )

```
import matplotlib.pyplot as plt
import numpy as np
x = np.arange(-8, 8, 0.1)
f = 1/(1 + np.exp(-x))
plt.plot(x, f)
plt.xlabel('x')
plt.ylabel('f(x)')
plt.show()
```



**Figure 4** The sigmoid function.



# The Sigmoid Function

- **Properties of Sigmoid Function**

- The sigmoid function returns a **real-valued** output.
- The first derivative of the sigmoid function will be non-negative or non-positive.
  - **Non-Negative:** If a number is greater than or equal to zero.
  - **Non-Positive:** If a number is less than or equal to Zero.

- **Sigmoid Function Usage**

- The Sigmoid function used for **binary classification** in the logistic regression model.
- While creating ANNs, uses the sigmoid function as the testing **activation function**.
- In statistics, the **sigmoid function** graphs are standard as a **cumulative distribution function**.

# Activation Function: tanh

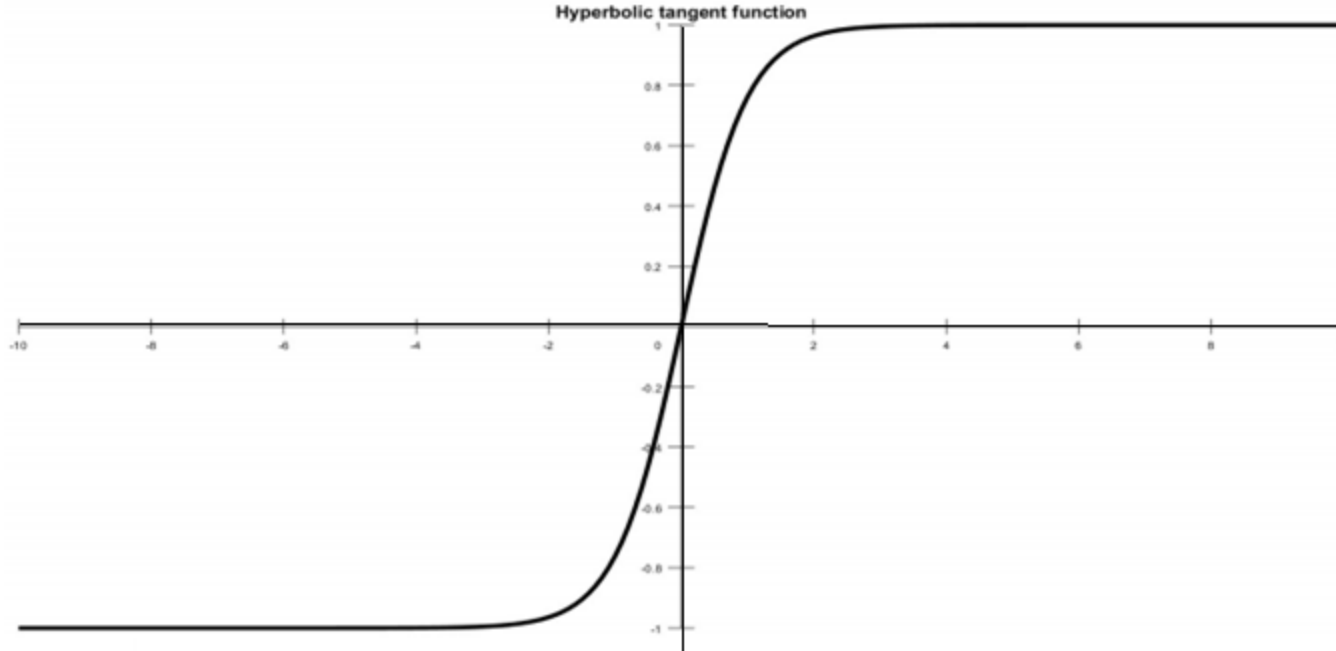
- Another popular NN activation function is the **tanh** (**hyperbolic tangent**) function.
- The **tanh** function is defined as:

$$f(x) = \tanh(x) = \frac{1 - e^{-2x}}{1 + e^{-2x}}$$

- It looks very similar to **sigmoid** function; in fact, **tanh function is a scaled sigmoid function**.
- As **sigmoid function**, this is also a **nonlinear function**, defined in the range of values **(-1, 1)**.
- The **gradient** (or **slope**) is stronger for **tanh** than **sigmoid** (the derivatives of its *exponential components* are more steep).

# Activation Function: tanh

- Deciding between **sigmoid** and **tanh** will depend on gradient strength requirement of an application. Like the **sigmoid**, **tanh** also has the **missing slope** problem. **Figure5** shows the **tanh** activation function:



It looks very similar to sigmoid function; in fact, it is a **scaled sigmoid** function.

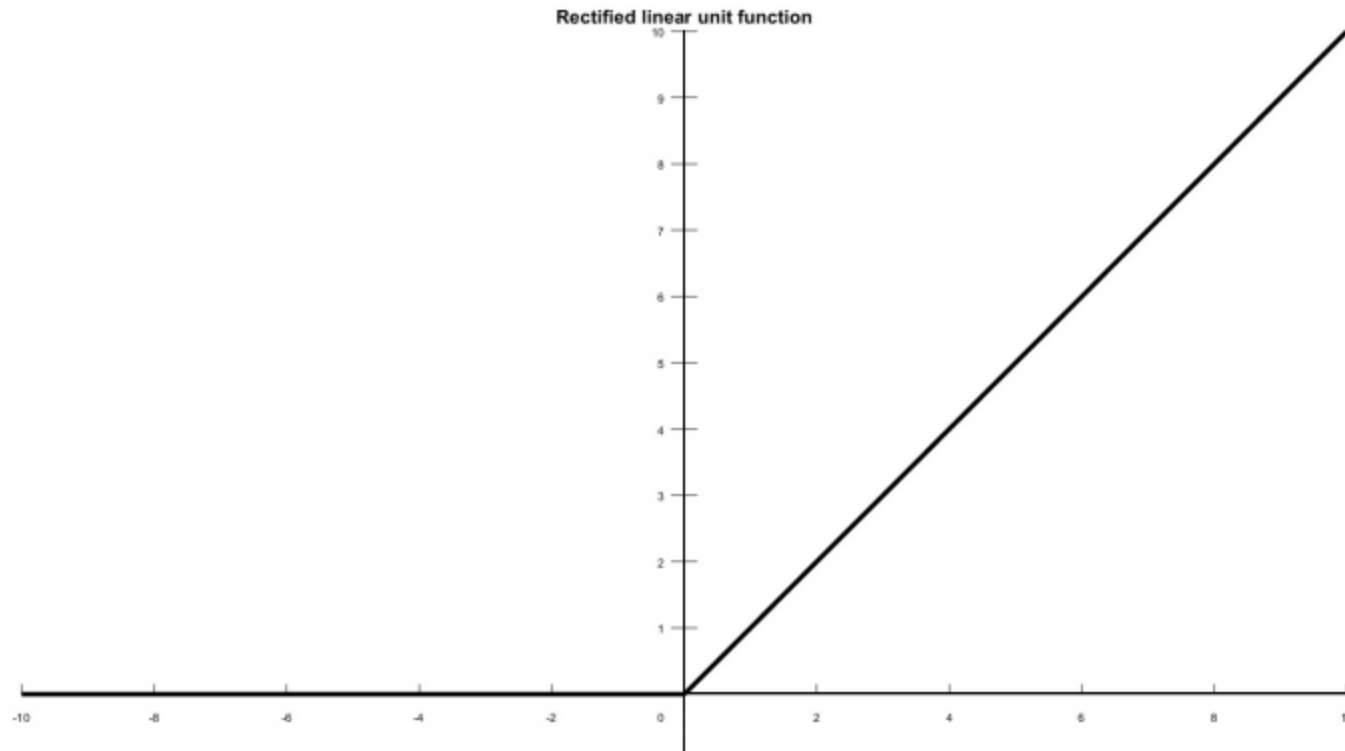
**Figure 5**

# Activation Function: ReLU

- **Rectified Linear Unit (ReLU)** is the most used activation function for applications based on **CNN** (Convolutional NN).
- It is a simple condition and has advantages over the other functions.
- The function is defined by the following formula:
$$f(x) = \begin{cases} 0 & \text{when } (x < 0) \\ x & \text{when } (x \geq 0) \end{cases}$$
- $f(x) = \max(x, 0)$
- The range of output is between **0** and **infinity**.
- **ReLU** finds applications in **computer vision** and **speech recognition**.

# Activation Function: ReLU

- **Figure6** shows a **ReLU** activation function:



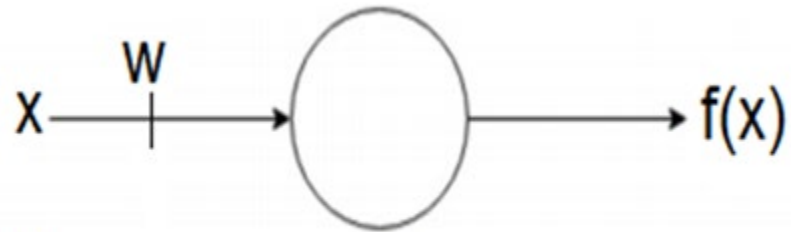
**Figure 6**

# Which activation functions to use?

- The **sigmoid** is the most used activation function, but it suffers from the following setbacks:
  - Since it uses a logistic model, the computations are time-consuming and complex.
  - It causes gradients to vanish, and no signals pass through the neurons at some point.
  - It is slow in convergence.
  - It is not zero-centered.

# Effects of Adjusting Weights

- Let's take a neuron node with only **one input** and **one output** (without a **bias** value) which is shown in **Figure 7**:

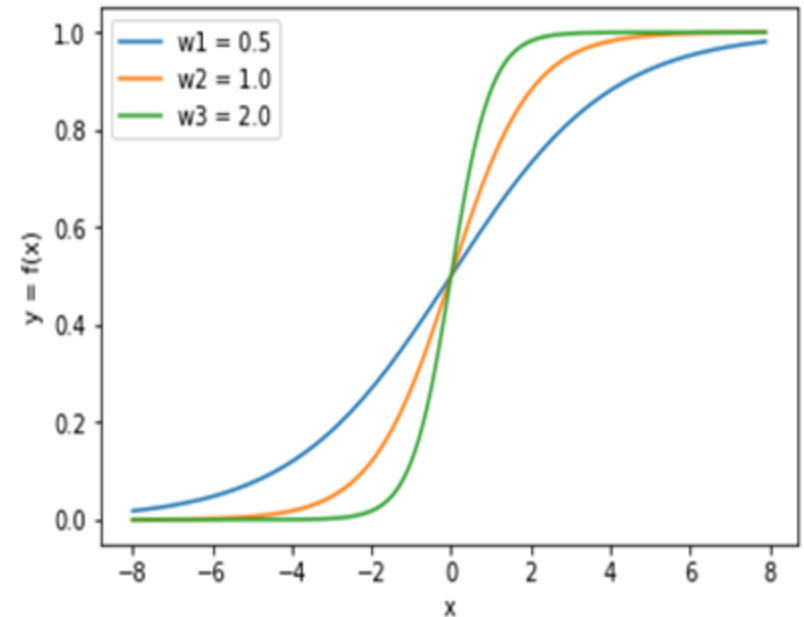


**Figure 7:** A single neuron node.

- The neuron's activation function, in this case, is the **sigmoid function**. What does changing **weight  $w$**  do in this simple network?
  - Figure 8** shows that changing the **weight** changes the slope of the output of the sigmoid activation function, which is helpful if we want to model different strengths of relationships between the **input** and **output** variables.

# Effects of Adjusting Weights

```
import matplotlib.pyplot as plt
import numpy as np
x = np.arange(-8, 8, 0.1)
w1 = 0.5
w2 = 1.0
w3 = 2.0
l1 = 'w1 = 0.5'
l2 = 'w2 = 1.0'
l3 = 'w3 = 2.0'
for w,l in [(w1, l1), (w2, l2), (w3, l3)]:
    f = 1/(1+np.exp(-x*w))
    plt.plot(x, f, label = l)
plt.xlabel('x')
plt.ylabel('y = f(x)')
plt.legend(loc = 2)
plt.show()
```



**Figure 8:** Effect of adjusting weights

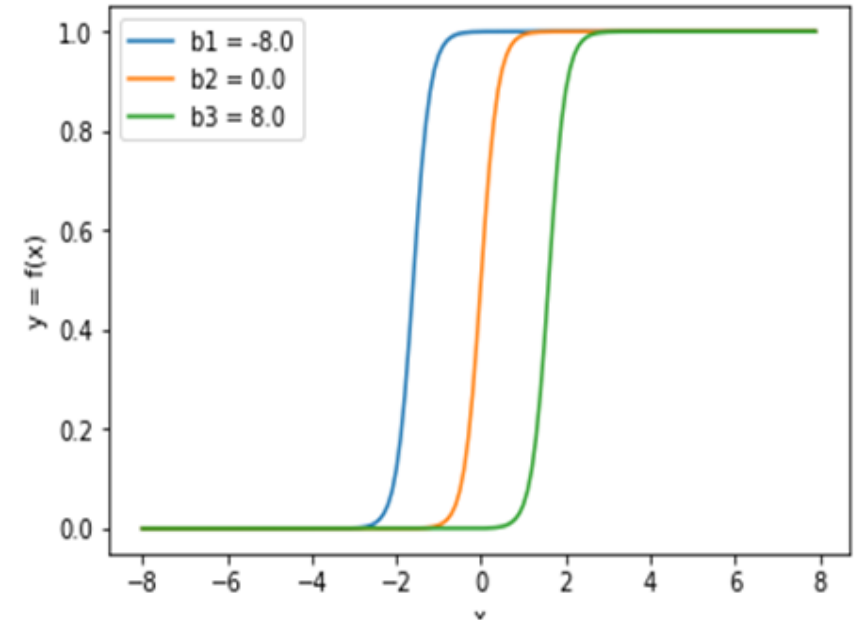


# Effects of Adjusting Bias

- The **weights** are real-valued numbers multiplied by the inputs and then summed up in the NN node.
  - In this case, the **w** has been increased to simulate a more defined “turn on” function.
- So, in other words, the **weighted input** to a neuron node with three inputs (**x<sub>1</sub>**, **x<sub>2</sub>**, and **x<sub>3</sub>**) and their respective weights (**w<sub>1</sub>**, **w<sub>2</sub>**, and **w<sub>3</sub>**) can be:  **$x_1w_1 + x_2w_2 + x_3w_3 + b$** 
  - where the **b** is the weight of the **+1** bias element of the neuron node, and
  - the inclusion of this **bias weight** enhances the flexibility of the neuron node.
  - The effect of **bias** adjustment is shown in **Figure 9**.

# Effects of Adjusting Bias

```
import matplotlib.pyplot as plt
import numpy as np
x = np.arange(-8, 8, 0.1)
w = 5.0
b1 = -8.0
b2 = 0.0
b3 = 8.0
l1 = 'b1 = -8.0'
l2 = 'b2 = 0.0'
l3 = 'b3 = 8.0'
for b, l in [(b1, l1), (b2, l2), (b3, l3)]:
    f = 1/(1 + np.exp(-(x*w) + b))
    plt.plot(x, f, label = l)
plt.xlabel('x')
plt.ylabel('y = f(x)')
plt.legend(loc = 2)
plt.show()
```



**Figure 9:** Effects of bias adjustments

# Effects of Adjusting Bias

- **Figure 9** shows that by varying the **bias weight  $b$** , you can change the **output** when the neuron node activates.
- Therefore, by adding a **bias** term, you can make a neuron node simulate a generic **if** function;
  - i.e. **if ( $x > z$ ) then 1 else 0.**
  - Without a **bias** term, you cannot vary the  **$z$**  of the **if statement**; it will always be stuck around 0.
- This is very useful if you are trying to simulate conditional relationships.

# Gradient Descent Optimization

- **Gradient descent** is an iterative method to find the *minimum* of a function (*gradient descent will be very clear from the **perceptron learning** section*).
- **Gradient descent** starts with an initial set of **weights**; in each step, it **decreases** each **weight** in proportion to its partial derivative:

$$w_i := w_i + \eta \times \frac{\partial \text{Error}_E(\bar{w})}{\partial w_i}$$

- where  $\eta$ , the gradient descent step size, is called the **learning rate**.
- **The learning rate**, as well as the features and the data, is given as input to the learning algorithm.
- The partial derivative specifies **how much a small change in the weight would change the error**.

# Gradient Descent Optimization

- Consider minimizing the **sum-of-squares error**. The error is the sum of all the examples. The partial derivative of a sum is the sum of the partial.
- For each example  $\mathbf{e}$ , let  $\delta = \text{val}(\mathbf{e}, \mathbf{Y}) - \text{pval}^w(\mathbf{e}, \mathbf{Y}) = \text{error}$ .
- Thus, each example  $\mathbf{e}$  updates each weight  $w_i$ :

$$w_i \quad := \quad w_i + \eta \times \delta \times \text{val}(e, X_i)$$

- Algorithm, ***LinearLearner*** ( $X, Y, E, \eta$ ), for learning a **linear function for minimizing the sum-of-square error (SOSE)**.

# Gradient Descent Optimization

- **Gradient descent** is an iterative method *to find the minimum of a function*.
- **Gradient descent** for *minimizing error* starts with an initial set of **weights**; in each step, *it decreases each weight* in proportion to its **partial derivative** (a *derivative* shows the sensitivity of change of a function's output with respect to its input.):

$$w_i := w_i - \eta * \frac{\partial}{\partial w_i} \text{error}(Es, \bar{w})$$

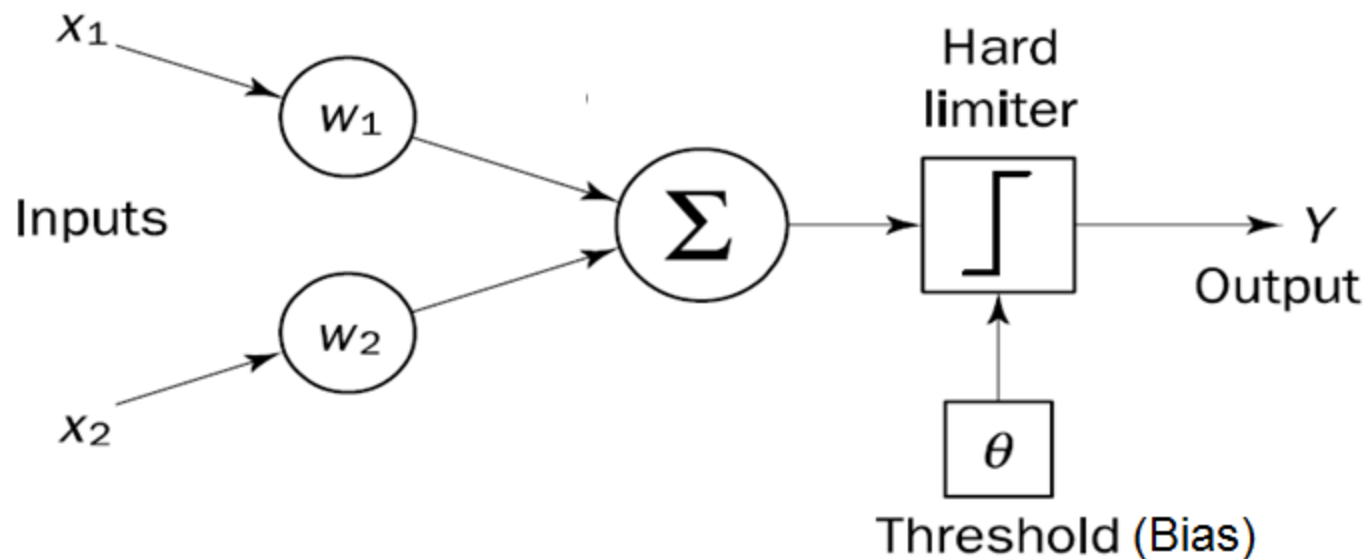
where  $\eta$ , the gradient descent step size, is called the **learning rate**. The learning rate, as well as the features and the data, is given as input to the learning algorithm. The partial derivative specifies how much a small change in the weight would change the error.

# Gradient Descent Optimization

- Neural learning keeps changing the **weights** until there is the greatest ***error reduction*** by an amount known as the **learning rate** ( $\alpha$ ).
  - **Learning rate** is a scalar parameter used to set the rate of adjustments to reduce the errors faster.
  - It adjusts **weights** and **biases** in a backpropagation learning process.
- **The higher the learning rate, the faster the algorithm will reduce the errors and the faster the training process.**

# The Perceptron

- The **perceptron** is the simplest form of neural net with **one neuron**.
- It consists of single **neuron** with adjustable *weights* and a **hard limiter** output function (depends on the application)
- A **perceptron** with two input is shown in **Figure 10**:

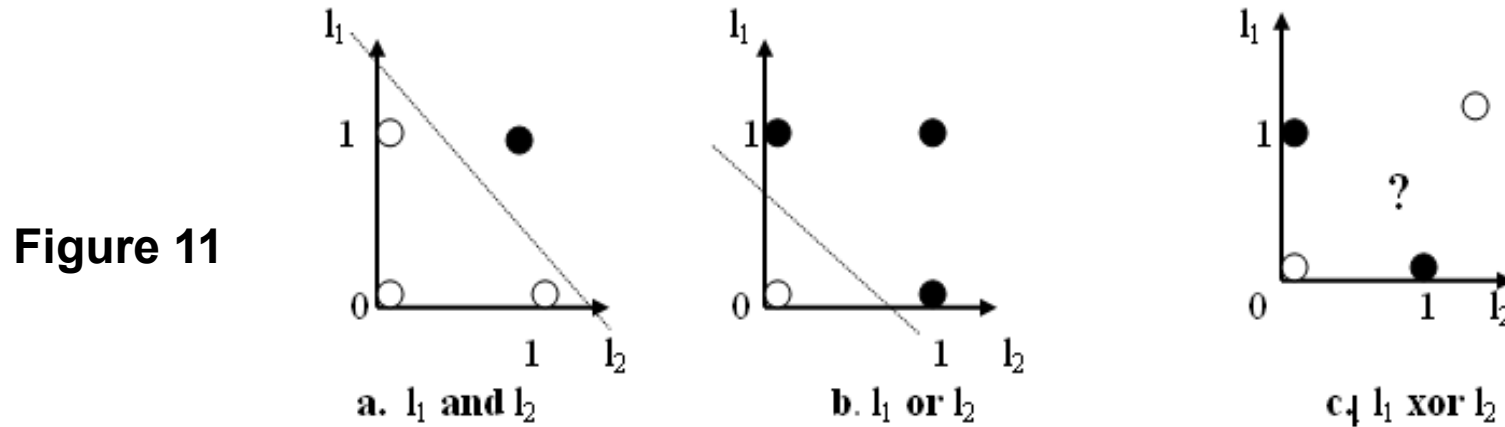


**Figure 10:** A perceptron with two input.



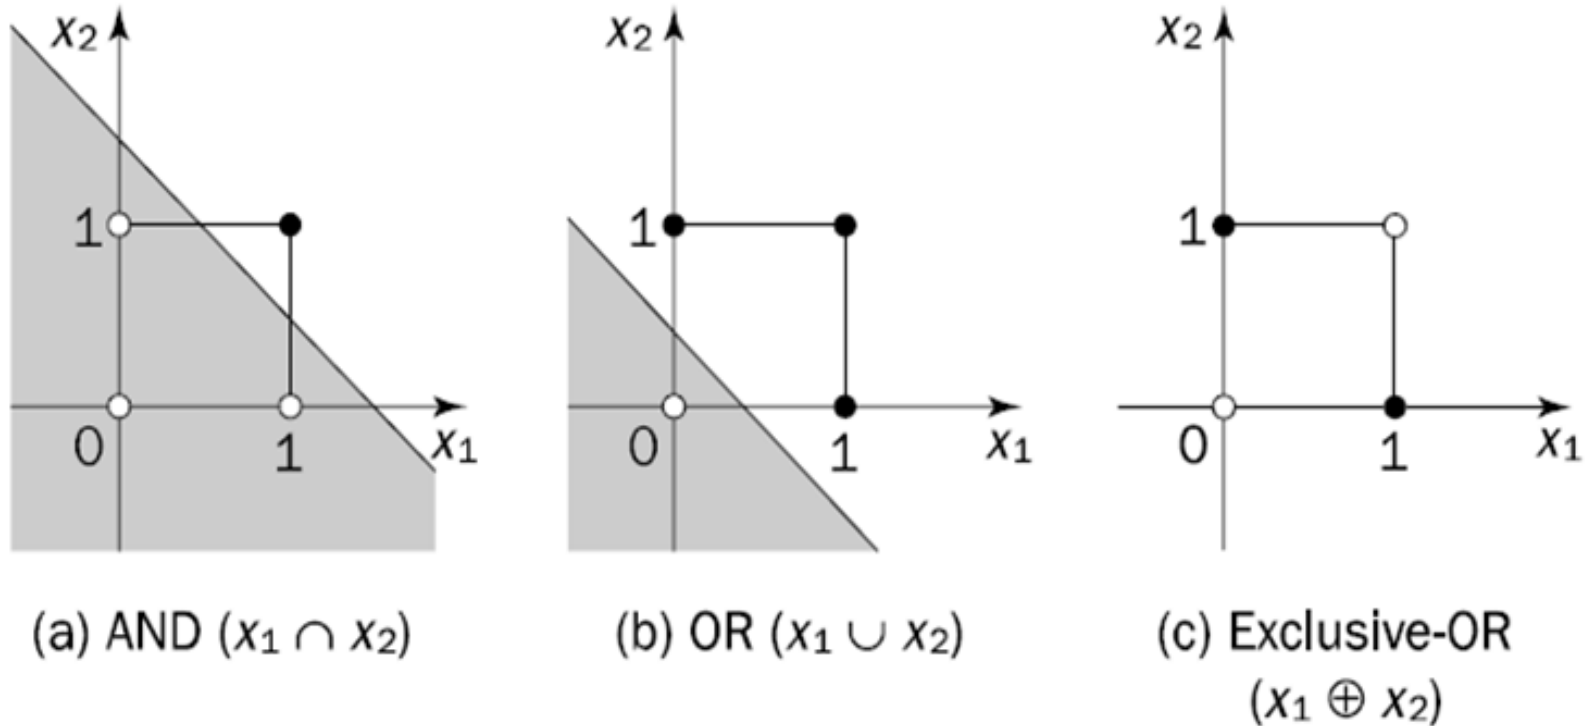
# Linear Separability in Perceptrons

- **Figure 11** and **Figure 12** show three different Boolean functions of two inputs, the AND, OR, and XOR functions



- Each function is represented as a 2D plot, based on the values of the two inputs (**black dots** indicate **1**, and **white dots** indicate **0**)
- A perceptron can represent a function only if there is some line that separates all the **white dots** from the **black dots** called a **linearly separable** function.
- Thus, a perceptron can represent **AND**, and **OR**, but **not XOR** !

# Linear Separability in Perceptrons



**Figure 12:** Two-dimensional plots of basic logical operations

# The Perceptron Learning Rule

- Rosenblatt first proposed the perceptron learning rule in 1960
- Using this rule, we can derive the **perceptron training algorithm** for classification tasks
- This is done by making small adjustments in the weights to reduce the difference between the **perceptron's actual (calculated/predicted) and desired (given/expected) outputs.**
  - The **initial weights** are randomly assigned in the range  $[-0.5, 0.5]$  and then updated to obtain the output consistent with the training examples

# The Perceptron Learning Rule

- If at iteration  $p$ , the **calculated/predicted** output is  $Y(p)$  and the **desired** output is  $Y_d(p)$ , then the **error**  $e$  at iteration  $p$  is given by
$$e(p) = Y_d(p) - Y(p), \text{ where } p = 1, 2, 3, \dots \quad (3)$$
- If the error  $e(p)$  is **positive**, we need to **increase**  $Y(p)$
- If the error is **negative**, we need to **decrease**  $Y(p)$
- Taking into the account that each perceptron input contributes  $x_i(p) \times w_i(p)$  with a **base** (threshold)  $\theta$  to the total input  $X(p)$
- Based on this concept, the **perceptron learning rule** is established:

$$w_i(p+1) = w_i(p) + \alpha \times x_i(p) \times e(p) \quad (4)$$

Where  $\alpha$  is the **learning rate**, a positive constant **less than the unity**

# Steps Behind Perceptron Learning Process

- Step1: Initialization

- Set initial weights  $w_1, w_2, \dots, w_n$  and base  $\theta$  to random numbers in the range  $[-0.5, 0.5]$

- Step2: Activation

- Activate the perceptron by applying inputs  $x_1(p), x_2(p), \dots, x_n(p)$  and **desired** output  $Y_d(p)$ . Calculate the **actual output**  $Y(p)$  at  $p = 1$ :

$$Y(p) = \text{step} \left[ \sum_{i=1}^n x_i(p) * w_i(p) - \theta \right] \quad (5)$$

- Where  $n$  is the number of the perceptron inputs and **step** is the **activation function**

# Steps Behind Perceptron Learning Process

- Step3: Weight training

- Update the weights of the perceptron

$$w_i(p+1) = w_i(p) + \Delta w_i(p) \quad (6)$$

- Where  $\Delta w_i(p)$  is the weight correction at iteration  $p$ . The weight correction is computed by the **delta rule**:

$$\Delta w_i(p) = \alpha \times x_i(p) \times e(p) \quad (7)$$

- Step4: Iteration

- Increase iteration  $p$  by one, go back to step 2 and repeat the process until **convergence** (all the  $y(p)$  focus with  $y_d(p)$  without error)

# Train Perceptron for AND and OR functions

- The truth tables for operations AND, OR and XOR are shown in **Table 1**. The perceptron must be trained to classify the input patterns

**Table 1** Truth tables for the basic logical operations

Input variables		AND	OR	Exclusive-OR
$x_1$	$x_2$	$x_1 \cap x_2$	$x_1 \cup x_2$	$x_1 \oplus x_2$
0	0	0	0	0
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0

- The training process for AND function is shown in **Table 2**

**Table 2** Example of perceptron learning: the logical operation AND

Epoch	$x_1$	$x_2$	$Y_d$	$w_1$	$w_2$	$Y$	$e$	$w_1$	$w_2$
1	0	0	0	0.3	-0.1	0	0	0.3	-0.1
	0	1	0	0.3	-0.1	0	0	0.3	-0.1
	1	0	0	0.3	-0.1	1	-1	0.2	-0.1
	1	1	1	0.2	-0.1	0	1	0.3	0.0
2	0	0	0	0.3	0.0	0	0	0.3	0.0
	0	1	0	0.3	0.0	0	0	0.3	0.0
	1	0	0	0.3	0.0	1	-1	0.2	0.0
	1	1	1	0.2	0.0	1	0	0.2	0.0
3	0	0	0	0.2	0.0	0	0	0.2	0.0
	0	1	0	0.2	0.0	0	0	0.2	0.0
	1	0	0	0.2	0.0	1	-1	0.1	0.0
	1	1	1	0.1	0.0	0	1	0.2	0.1
4	0	0	0	0.2	0.1	0	0	0.2	0.1
	0	1	0	0.2	0.1	0	0	0.2	0.1
	1	0	0	0.2	0.1	1	-1	0.1	0.1
	1	1	1	0.1	0.1	1	0	0.1	0.1
5	0	0	0	0.1	0.1	0	0	0.1	0.1
	0	1	0	0.1	0.1	0	0	0.1	0.1
	1	0	0	0.1	0.1	0	0	0.1	0.1
	1	1	1	0.1	0.1	1	0	0.1	0.1

Threshold:  $\theta = 0.2$ ; learning rate:  $\alpha = 0.1$ .



# Why Can a Perceptron Learn Only Linearly Separable Functions?

- The fact that a perceptron can learn only **linearly separable functions** based on the following equation:

$$X = \sum_{i=1}^n x_i w_i$$
$$Y = \begin{cases} +1 & \text{if } X \geq \theta \\ -1 & \text{if } X < \theta \end{cases}$$

- The perceptron output  $Y$  is 1 only if the total weighted input  $X$  is greater than or equal to the threshold,  $\theta$ . This means that the entire input space is divided in two along a boundary defined by  $X = \theta$
- A separating line for the operation AND is defined by the equation  $x_1 w_1 + x_2 w_2 = \theta$

# Why Can a Perceptron Learn Only Linearly Separable Functions?

- If we substitute values for weights  $w_1$  and  $w_2$  and threshold  $\theta$  given in **Table 2**, we obtain one of the possible separating lines as (see 5<sup>th</sup> iteration):

$$0.1x_1 + 0.1x_2 = 0.2 \quad \text{or} \quad x_1 + x_2 = 2$$

- Thus, the region below the boundary line, where the output is **0**, is given by  $x_1 + x_2 - 2 < 0$ ,
- And the region above this line, where the output is **1**, is given by  $x_1 + x_2 - 2 \geq 0$
- So a perceptron can learn only **linear separable functions** and there are not many such functions!

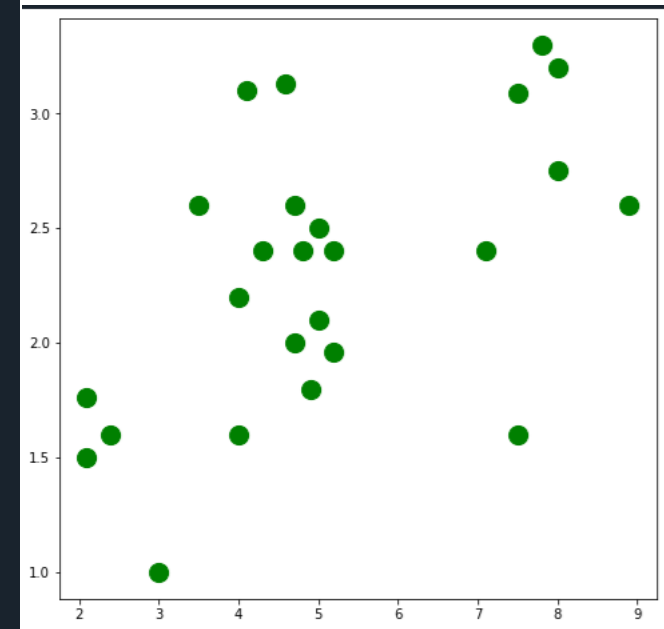
# Linear Regression with Perceptron

<https://www.analyticsvidhya.com/blog/2021/11/a-comprehensive-guide-to-linear-regression-with-pytorch/>

```
import numpy as np
import matplotlib.pyplot as plt
import torch
from torch.autograd import Variable
#Dataset is created using NumPy arrays.
x_train = np.array ([[4.7], [2.4], [7.5], [7.1], [4.3],
                     [7.8], [8.9], [5.2], [4.59], [2.1],
                     [8], [5], [7.5], [5], [4],
                     [8], [5.2], [4.9], [3], [4.7],
                     [4], [4.8], [3.5], [2.1], [4.1]],
                    dtype = np.float32)

y_train = np.array ([[2.6], [1.6], [3.09], [2.4], [2.4],
                     [3.3], [2.6], [1.96], [3.13], [1.76],
                     [3.2], [2.1], [1.6], [2.5], [2.2],
                     [2.75], [2.4], [1.8], [1], [2],
                     [1.6], [2.4], [2.6], [1.5], [3.1]],
                    dtype = np.float32)

#Visualizing the data.
plt.figure(figsize=(8,8))
plt.scatter(x_train, y_train, c='green', s=200, label='Original data')
plt.show()
```



# Linear Regression with Perceptron

<https://www.analyticsvidhya.com/blog/2021/11/a-comprehensive-guide-to-linear-regression-with-pytorch/>

```
7 import numpy as np
8 import matplotlib.pyplot as plt
9 import torch
10 from torch.autograd import Variable
11 #Dataset is created using NumPy arrays.
12 x_train = np.array ([[4.7], [2.4], [7.5], [7.1], [4.3],
13                     [7.8], [8.9], [5.2], [4.59], [2.1],
14                     [8], [5], [7.5], [5], [4],
15                     [8], [5.2], [4.9], [3], [4.7],
16                     [4], [4.8], [3.5], [2.1], [4.1]],
17                     dtype = np.float32)
18
19 y_train = np.array ([[2.6], [1.6], [3.09], [2.4], [2.4],
20                     [3.3], [2.6], [1.96], [3.13], [1.76],
21                     [3.2], [2.1], [1.6], [2.5], [2.2],
22                     [2.75], [2.4], [1.8], [1], [2],
23                     [1.6], [2.4], [2.6], [1.5], [3.1]],
24                     dtype = np.float32)
25
26 #Visualizing the data.
27 plt.figure(figsize=(8,8))
28 plt.scatter(x_train, y_train, c='green', s=200, label='Original data')
29 plt.show()
30
31 X_train = torch.from_numpy(x_train) #Convert numpy arrays into Pytorch arrays
32 Y_train = torch.from_numpy(y_train) #Convert numpy arrays into Pytorch arrays
33 print('requires_grad for X_train: ', X_train.requires_grad)
34 print('requires_grad for Y_train: ', Y_train.requires_grad)
```

X_train	Y_train
tensor([[4.7000],	tensor([[2.6000],
[2.4000],	[1.6000],
[7.5000],	[3.0900],
[7.1000],	[2.4000],
[4.3000],	[2.4000],
[7.8000],	[3.3000],
[8.9000],	[2.6000],
[5.2000],	[1.9600],
[4.5900],	[3.1300],
[2.1000],	[1.7600],
[8.0000],	[3.2000],
[5.0000],	[2.1000],
[7.5000],	[1.6000],
[5.0000],	[2.5000],
[4.0000],	[2.2000],
[8.0000],	[2.7500],
[5.2000],	[2.4000],
[4.9000],	[1.8000],
[3.0000],	[1.0000],
[4.7000],	[2.0000],
[4.0000],	[1.6000],
[4.8000],	[2.4000],
[3.5000],	[2.6000],
[2.1000],	[1.5000],
[4.1000]]	[3.1000]])

# Linear Regression with Perceptron

<https://www.analyticsvidhya.com/blog/2021/11/a-comprehensive-guide-to-linear-regression-with-pytorch/>

```
39 #Parameters of the model are W1 and b1, which is weight and bias respectively.
40 #Let us have a look at the parameters that are defined.
41 input_size = 1
42 hidden_size = 1
43 output_size = 1
44 learning_rate = 0.001
45 w1 = torch.rand(input_size,
46                 hidden_size,
47                 requires_grad=True)
48
49 b1 = torch.rand(hidden_size,
50                 output_size,
51                 requires_grad=True)
52
53 for iter in range(1, 5000):
54     y_pred = X_train.mm(w1).clamp(min=0).add(b1)
55     loss = (y_pred - Y_train).pow(2).sum()
56     if iter % 100 == 0:
57         print(iter, loss.item())
58     loss.backward()
59     with torch.no_grad():
60         w1 -= learning_rate * w1.grad
61         b1 -= learning_rate * b1.grad
62         w1.grad.zero_()
63         b1.grad.zero_()
64 #Let us check the optimized value for W1 and b1:
65 print(w1, b1)
```

2700	loss =	6.119896411895752
2800	loss =	6.119896411895752
2900	loss =	6.119896411895752
3000	loss =	6.119896411895752
3100	loss =	6.119896411895752
3200	loss =	6.119896411895752
3300	loss =	6.119896411895752
3400	loss =	6.119896411895752
3500	loss =	6.119896411895752
3600	loss =	6.119896411895752
3700	loss =	6.119896411895752
3800	loss =	6.119896411895752
3900	loss =	6.119896411895752
4000	loss =	6.119896411895752
4100	loss =	6.119896411895752
4200	loss =	6.119896411895752
4300	loss =	6.119896411895752
4400	loss =	6.119896411895752
4500	loss =	6.119896411895752
4600	loss =	6.119896411895752
4700	loss =	6.119896411895752
4800	loss =	6.119896411895752
4900	loss =	6.119896411895752

# Linear Regression with Perceptron

<https://www.analyticsvidhya.com/blog/2021/11/a-comprehensive-guide-to-linear-regression-with-pytorch/>

```
68 #Let us check the optimized value for W1 and b1:
69 print ('w1: ', w1)
70 print ('b1: ', b1)
71 #Getting the prediction values using the weights in the linear equation.
72 getX = float(input("Enter x_tain value: "))
73 predicted_in_tensor = getX * w1 + b1
74 print ("The predicted output = ", Variable(predicted_in_tensor) )
```

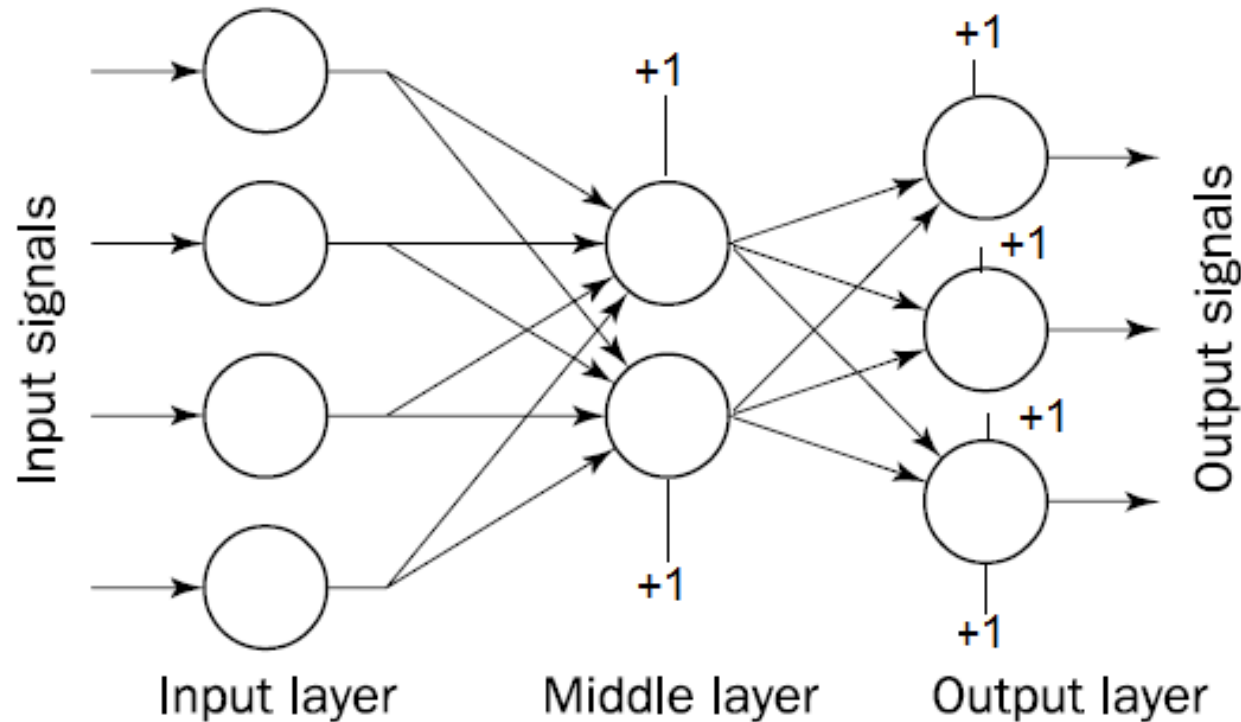
```
w1:  tensor([[0.1751]], requires_grad=True)
b1:  tensor([[1.4045]], requires_grad=True)
Enter x_tain value: 2.4
The predicted output =  tensor([[1.8247]])
```

# Perceptron Exercises

1. Create two input AND functions and two input OR functions from perceptron ( assume suitable weights and threshold value)
2. Train a perceptron for getting 2-input OR function. Assume  $w_1 = 0.3$ ,  $w_2 = -0.1$ ,  $\theta = 0.2$  and  $\alpha = 0.1$
3. Consider a perceptron with two real-valued inputs and an output unit with a *sigmoid activation* function. All the initial weights and the bias (threshold) equal 0.5. Assume that the output should be 1 for the input  $x_1 = 0.7$  and  $x_2 = -0.6$ . Show how the delta rule supports the training of the neuron (assume  $\alpha = 0.1$ )
4. Predict the value of **BMI** (Body Mass Index) from the ***Height*** and ***Weight*** of a person using a **perceptron**. Use the ***bmi.csv*** dataset for training (ignore its '***gender***' column). **Don't use any Python perceptron library**; use the perceptron learning steps from this lecture slide (use the *sigmoid activation* function and modify your dataset accordingly).

# Structure of a Multilayer ANN

- The structure of an **Multilayer ANN** is shown in **Figure 13**:



**Figure 13:** The structure of an ANN